

fc_project_li (copy)

August 15, 2021

```
[51]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import plotly
import plotly.express as px
import plotly.graph_objs as go
from plotly.offline import init_notebook_mode, plot, iplot
import cv2
import random
import os
import glob
import tensorflow.keras
import seaborn as sns
from tensorflow.keras.applications.resnet50 import ResNet50
from keras_applications import xception
from tqdm.notebook import tqdm
import albumentations as A
import tensorflow as tf
from tensorflow.keras.layers import Conv2D, Flatten, MaxPooling2D, Dense,
↳Dropout , BatchNormalization
from tensorflow.keras.models import Sequential
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.applications.vgg16 import VGG16
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report, confusion_matrix
from sklearn.metrics import roc_curve, roc_auc_score, auc

from sklearn.utils.multiclass import unique_labels
from sklearn.metrics import accuracy_score, precision_score, recall_score,
↳f1_score, fbeta_score, classification_report
from sklearn.metrics import roc_curve, precision_recall_curve, roc_auc_score
from tensorflow.keras.callbacks import ReduceLROnPlateau
import torch
import warnings
warnings.filterwarnings('ignore')
warnings.simplefilter('ignore')
import time
```

```
[178]: device=torch.device('cuda')
print('Num GPUs available', len(tf.config.experimental.
      ↪list_physical_devices('GPU')))
```

Num GPUs available 1

```
[53]: cwd=os.getcwd()
wd=cwd+'/chest_xray'
```

```
[54]: train_data=glob.glob(wd+'/train/**/*.jpeg')
test_data=glob.glob(wd+'/test/**/*.jpeg')
val_data=glob.glob(wd+'/val/**/*.jpeg')
```

```
[55]: print("//"*20)
print(f"Training Set has: {len(train_data)} images")
print(f"Testing Set has: {len(test_data)} images")
print(f"Validation Set has: {len(val_data)} images")
print("//"*20)
```


```
////////////////////////////////////
Training Set has: 5216 images
Testing Set has: 624 images
Validation Set has: 16 images
////////////////////////////////////
```

```
[56]: sets = ["train", "test", "val"]
all_pneumonia = []
all_normal = []
results={}
results['model']=[]
results['total_param']=[]
results['time']=[]
results['score']=[]

for cat in sets:
    path = os.path.join(wd, cat)
    norm = glob.glob(os.path.join(path, "NORMAL/*.jpeg"))
    pneu = glob.glob(os.path.join(path, "PNEUMONIA/*.jpeg"))
    all_normal.extend(norm)
    all_pneumonia.extend(pneu)
print(" "*20)
print(f"Total Pneumonia Images: {len(all_pneumonia)}")
print(f"Total Normal Images: {len(all_normal)}")
print(" "*20)
```


Total Pneumonia Images: 4273
Total Normal Images: 1583

```
[57]: ## Plotly chart for Class distribution
labels = ["Normal", 'Pneumonia ']
values = [len(all_normal), len(all_pneumonia)]
colors = ['red', 'blue']
fig = go.Figure(data=[go.Pie(labels=labels,
                             values=values)])
fig.update_traces(hoverinfo='value', textinfo='label+percent', textfont_size=20,
                  marker=dict(colors=colors, line=dict(color='#000000',
→width=3)))
fig.update_layout(title="Image Category Distribution", title_x=0.5)
iplot(fig)
```



```
[58]: random.shuffle(all_normal)
random.shuffle(all_pneumonia)
images = all_normal[:11] + all_pneumonia[:10]
```

```
[59]: fig=plt.figure(figsize=(15, 10))
columns = 4; rows = 5
for i in range(1, columns*rows+1):
    img = cv2.imread(images[i])
    img = cv2.resize(img, (128, 128))
    fig.add_subplot(rows, columns, i)
    plt.imshow(img)
    plt.axis(False)
```



1 Base Model

```
[60]: #data augmentation
train_gen = ImageDataGenerator(
    rescale=1/255.,
    featurewise_center=False, # set input mean to 0 over the dataset
    samplewise_center=False, # set each sample mean to 0
    featurewise_std_normalization=False, # divide inputs by std of the
    ↪ dataset
    samplewise_std_normalization=False, # divide each input by its std
    zca_whitening=False, # apply ZCA whitening
    rotation_range = 30, # randomly rotate images in the range (degrees, 0
    ↪ to 180)
    zoom_range = 0.4, # Randomly zoom image
    width_shift_range=0.1, # randomly shift images horizontally (fraction
    ↪ of total width)
    height_shift_range=0.1, # randomly shift images vertically (fraction
    ↪ of total height)
    horizontal_flip = True, # randomly flip images
    vertical_flip=False,
```

```

        validation_split=0.2) #validation size
test_gen=ImageDataGenerator(
    rescale=1/255.,
    featurewise_center=False, # set input mean to 0 over the dataset
    samplewise_center=False, # set each sample mean to 0
    featurewise_std_normalization=False, # divide inputs by std of the
    ↪ dataset
    samplewise_std_normalization=False, # divide each input by its std
    zca_whitening=False, # apply ZCA whitening
    horizontal_flip = False, # randomly flip images
    vertical_flip=False) # randomly flip images

```

```

[77]: batch_size=64
nb_epochs=90
train_generator = train_gen.flow_from_directory(
    wd+"/train",
    batch_size=batch_size,
    target_size=(224, 224), #class_mode="binary"
    # class_mode='binary',
    subset='training'
)

validation_generator = train_gen.flow_from_directory(
    wd+"/train", # same directory as training data
    batch_size=batch_size,
    target_size=(224, 224),
    # class_mode='binary',
    subset='validation') # set as validation data
test_generator=test_gen.flow_from_directory(
    wd+"/test", # same directory as training data
    batch_size=1,
    target_size=(224, 224),
    shuffle=False)

```

Found 4173 images belonging to 2 classes.

Found 1043 images belonging to 2 classes.

Found 624 images belonging to 2 classes.

```

[62]: model = Sequential()
model.add(Conv2D(32,(3,3),strides=(1, 1),activation='relu',padding='same',
    ↪ input_shape=(224, 224, 3)))
model.add(MaxPooling2D(pool_size=(2,2)))
model.add(Conv2D(64,(3,3),strides=(1, 1) ,padding='same',activation='relu'))
model.add(MaxPooling2D(pool_size=(2,2)))
model.add(Conv2D(128,(3,3),strides=(1, 1),padding='same', activation='relu'))
model.add(MaxPooling2D(pool_size=(2,2)))
model.add(Conv2D(256,(3,3),strides=(1, 1),padding='same', activation='relu'))

```

```

model.add(MaxPooling2D(pool_size=(2,2)))
model.add(Flatten())

model.add(Dense(128, activation='relu'))
model.add(Dense(64, activation='relu'))
model.add(Dense(2, activation='softmax'))
model.summary()

```

Model: "sequential_4"

Layer (type)	Output Shape	Param #
conv2d_16 (Conv2D)	(None, 224, 224, 32)	896
max_pooling2d_16 (MaxPooling)	(None, 112, 112, 32)	0
conv2d_17 (Conv2D)	(None, 112, 112, 64)	18496
max_pooling2d_17 (MaxPooling)	(None, 56, 56, 64)	0
conv2d_18 (Conv2D)	(None, 56, 56, 128)	73856
max_pooling2d_18 (MaxPooling)	(None, 28, 28, 128)	0
conv2d_19 (Conv2D)	(None, 28, 28, 256)	295168
max_pooling2d_19 (MaxPooling)	(None, 14, 14, 256)	0
flatten_4 (Flatten)	(None, 50176)	0
dense_12 (Dense)	(None, 128)	6422656
dense_13 (Dense)	(None, 64)	8256
dense_14 (Dense)	(None, 2)	130
Total params: 6,819,458		
Trainable params: 6,819,458		
Non-trainable params: 0		

```

[64]: ## Compile the model
model.compile(optimizer='adam', loss='categorical_crossentropy',
             metrics=['accuracy'])

```

```

[65]: early_stopping_cb = tf.keras.callbacks.
      ↪ EarlyStopping(patience=5, restore_best_weights=True)

```

```

start=time.time()
history_base = model.fit_generator(
    train_generator,
    epochs = nb_epochs,
    validation_data=validation_generator,
    steps_per_epoch=train_generator.samples // batch_size,
    validation_steps = validation_generator.samples // batch_size,
    callbacks=[early_stopping_cb],
    verbose=1)
run_time=time.time()-start

```

Epoch 1/90

65/65 [=====] - 73s 1s/step - loss: 0.6386 - accuracy: 0.6670 - val_loss: 0.4465 - val_accuracy: 0.7695

Epoch 2/90

65/65 [=====] - 72s 1s/step - loss: 0.4128 - accuracy: 0.8038 - val_loss: 0.4406 - val_accuracy: 0.8076

Epoch 3/90

65/65 [=====] - 72s 1s/step - loss: 0.3380 - accuracy: 0.8420 - val_loss: 0.3415 - val_accuracy: 0.8281

Epoch 4/90

65/65 [=====] - 72s 1s/step - loss: 0.3384 - accuracy: 0.8502 - val_loss: 0.3706 - val_accuracy: 0.8125

Epoch 5/90

65/65 [=====] - 72s 1s/step - loss: 0.3088 - accuracy: 0.8628 - val_loss: 0.3360 - val_accuracy: 0.8369

Epoch 6/90

65/65 [=====] - 72s 1s/step - loss: 0.2778 - accuracy: 0.8751 - val_loss: 0.3408 - val_accuracy: 0.8281

Epoch 7/90

65/65 [=====] - 73s 1s/step - loss: 0.2877 - accuracy: 0.8818 - val_loss: 0.2947 - val_accuracy: 0.8721

Epoch 8/90

65/65 [=====] - 72s 1s/step - loss: 0.2526 - accuracy: 0.8893 - val_loss: 0.2901 - val_accuracy: 0.8750

Epoch 9/90

65/65 [=====] - 72s 1s/step - loss: 0.2605 - accuracy: 0.8899 - val_loss: 0.2696 - val_accuracy: 0.8740

Epoch 10/90

65/65 [=====] - 72s 1s/step - loss: 0.2233 - accuracy: 0.9071 - val_loss: 0.2904 - val_accuracy: 0.8643

Epoch 11/90

65/65 [=====] - 73s 1s/step - loss: 0.2166 - accuracy: 0.9127 - val_loss: 0.3759 - val_accuracy: 0.8672

Epoch 12/90

65/65 [=====] - 72s 1s/step - loss: 0.2131 - accuracy: 0.9134 - val_loss: 0.2718 - val_accuracy: 0.8701

Epoch 13/90
65/65 [=====] - 72s 1s/step - loss: 0.1961 - accuracy: 0.9254 - val_loss: 0.2658 - val_accuracy: 0.8896
Epoch 14/90
65/65 [=====] - 73s 1s/step - loss: 0.2063 - accuracy: 0.9152 - val_loss: 0.2518 - val_accuracy: 0.8887
Epoch 15/90
65/65 [=====] - 73s 1s/step - loss: 0.1735 - accuracy: 0.9290 - val_loss: 0.2071 - val_accuracy: 0.8994
Epoch 16/90
65/65 [=====] - 72s 1s/step - loss: 0.1903 - accuracy: 0.9255 - val_loss: 0.2388 - val_accuracy: 0.8916
Epoch 17/90
65/65 [=====] - 73s 1s/step - loss: 0.1674 - accuracy: 0.9362 - val_loss: 0.1923 - val_accuracy: 0.9189
Epoch 18/90
65/65 [=====] - 73s 1s/step - loss: 0.1703 - accuracy: 0.9323 - val_loss: 0.2189 - val_accuracy: 0.9189
Epoch 19/90
65/65 [=====] - 73s 1s/step - loss: 0.1608 - accuracy: 0.9299 - val_loss: 0.1756 - val_accuracy: 0.9248
Epoch 20/90
65/65 [=====] - 72s 1s/step - loss: 0.1528 - accuracy: 0.9384 - val_loss: 0.1793 - val_accuracy: 0.9229
Epoch 21/90
65/65 [=====] - 73s 1s/step - loss: 0.1474 - accuracy: 0.9442 - val_loss: 0.1797 - val_accuracy: 0.9297
Epoch 22/90
65/65 [=====] - 72s 1s/step - loss: 0.1444 - accuracy: 0.9414 - val_loss: 0.1848 - val_accuracy: 0.9219
Epoch 23/90
65/65 [=====] - 73s 1s/step - loss: 0.1515 - accuracy: 0.9418 - val_loss: 0.1706 - val_accuracy: 0.9209
Epoch 24/90
65/65 [=====] - 72s 1s/step - loss: 0.1417 - accuracy: 0.9494 - val_loss: 0.1822 - val_accuracy: 0.9238
Epoch 25/90
65/65 [=====] - 72s 1s/step - loss: 0.1363 - accuracy: 0.9457 - val_loss: 0.2172 - val_accuracy: 0.9131
Epoch 26/90
65/65 [=====] - 73s 1s/step - loss: 0.1395 - accuracy: 0.9458 - val_loss: 0.3130 - val_accuracy: 0.8750
Epoch 27/90
65/65 [=====] - 72s 1s/step - loss: 0.1778 - accuracy: 0.9325 - val_loss: 0.1891 - val_accuracy: 0.9131
Epoch 28/90
65/65 [=====] - 72s 1s/step - loss: 0.1400 - accuracy: 0.9406 - val_loss: 0.1916 - val_accuracy: 0.9092


```
[66]: score = model.evaluate_generator(test_generator, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])
```

Test score: 0.4663000702857971
Test accuracy: 0.807692289352417


```
[67]: #saving the model and histories
#tf.keras.models.save_model(filepath=cwd, model=model)
model.save(cwd+'/assets/model_base')
np.save(cwd+'/assets/history_base.npy', history_base.history)
results['model'].append('base_model')
results['total_param'].append(6819458)
results['time'].append(run_time)
results['score'].append(score[1])
```

INFO:tensorflow:Assets written to:
/home/jovyan/work/fc_project/assets/model_base/assets

```
[68]: fig, (ax1, ax2) = plt.subplots(1,2, figsize=(12,5))

ax1.plot(history_base.history['loss'])
ax1.plot(history_base.history['val_loss'])
ax1.set_title('loss')
ax1.set_ylabel('loss')
ax1.set_xlabel('epoch')
ax1.legend(['train tuned', 'test tuned', 'train', 'test'], loc='upper right')
ax2.plot(history_base.history['accuracy'])
ax2.plot(history_base.history['val_accuracy'])

ax2.set_title('accuracy')
ax2.set_ylabel('accuracy')
ax2.set_xlabel('epoch')
ax2.legend(['train tuned', 'test tuned', 'train', 'test'], loc='lower right')
fig.suptitle('loss and accuracy')
plt.show()
```



```
[78]: STEP_SIZE_TEST=624
test_generator.reset()
preds = model.predict(test_generator,
steps=STEP_SIZE_TEST,
verbose=1)
```

624/624 [=====] - 7s 7ms/step

```
[79]: #scale the output of the model to [0,1]
#scaled_preds=(preds[:,1]-min(preds[:,1]))/(max(preds[:,1])-min(preds[:,1]))
pos = [i for i, j in zip(preds[:,1], test_generator.classes) if j == 1]
neg = [i for i, j in zip(preds[:,1], test_generator.classes) if j == 0]

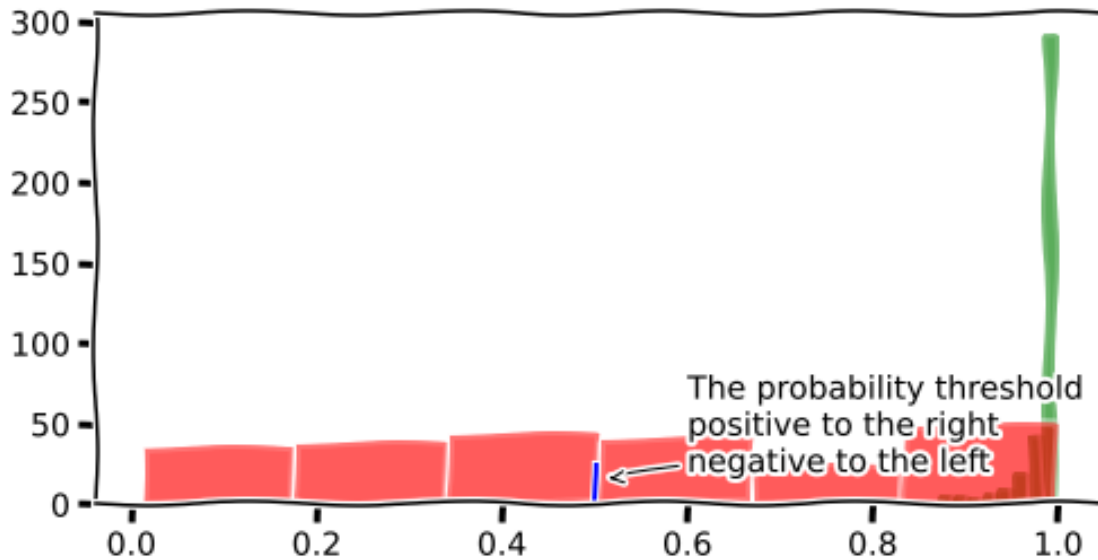
with plt.xkcd():
    fig = plt.figure(figsize=(8, 4))

    sns.distplot(pos, hist = True, kde = False, color='g',
                  kde_kws = {'shade': True, 'linewidth': 3})

    sns.distplot(neg, hist = True, kde = False, color='r',
                  kde_kws = {'shade': True, 'linewidth': 3})

    plt.plot([0.5, 0.5], [0, 25], '-b')
    plt.annotate(
        'The probability threshold\npositive to the right\nnegative to the
→left',
        xy=(0.51, 15), arrowprops=dict(arrowstyle='->'), xytext=(0.6, 20))

plt.show()
```



```
[80]: def plot_confusion_matrix(cm,
                                target_names,
                                title='Confusion matrix',
                                cmap=None,
                                normalize=True):

    """
    Given a scikit-learn confusion matrix (CM), make a nice plot.

    Arguments
    -----
    cm:                Confusion matrix from sklearn.metrics.confusion_matrix

    target_names:      Given classification classes, such as [0, 1, 2]
                        The class names, for example, ['high', 'medium', 'low']

    title:             The text to display at the top of the matrix

    cmap:             The gradient of the values displayed from matplotlib.pyplot.cm
                        See http://matplotlib.org/examples/color/colormaps_reference.
    ↪html              `plt.get_cmap('jet')` or `plt.cm.Blues`

    normalize:        If `False`, plot the raw numbers
                        If `True`, plot the proportions

    Usage
    -----
```

```

    plot_confusion_matrix(cm          = cm,                # Confusion matrix
    ↪matrix created by

                                # `sklearn.metrics.confusion_matrix`

                                normalize = True,           # Show proportions
                                target_names = y_labels_vals, # List of names of
    ↪of the classes

                                title      = best_estimator_name) # Title of graph

    Citation
    -----
    http://scikit-learn.org/stable/auto\_examples/model\_selection/plot\_confusion\_matrix.html

    """
    import matplotlib.pyplot as plt
    import numpy as np
    import itertools

    accuracy = np.trace(cm) / float(np.sum(cm))
    misclass = 1 - accuracy

    if cmap is None:
        cmap = plt.get_cmap('Blues')
    plt.figure(figsize=(8, 6))
    plt.imshow(cm, interpolation='nearest', cmap=cmap)
    plt.title(title)
    plt.colorbar()

    if target_names is not None:
        tick_marks = np.arange(len(target_names))
        plt.xticks(tick_marks, target_names, rotation=45)
        plt.yticks(tick_marks, target_names)

    if normalize:
        cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]

    thresh = cm.max() / 1.5 if normalize else cm.max() / 2
    for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
        if normalize:
            plt.text(j, i, "{:0.4f}".format(cm[i, j]),
                     horizontalalignment="center",
                     color="white" if cm[i, j] > thresh else "black")
        else:
            plt.text(j, i, "{:,}".format(cm[i, j]),
                     horizontalalignment="center",

```

```

        color="white" if cm[i, j] > thresh else "black")

plt.tight_layout()
plt.ylabel('True label')
plt.xlabel('Predicted label\naccuracy={:0.4f}; misclass={:0.4f}'.
    ↪format(accuracy, misclass))
plt.show()

```

```

[81]: confusion = confusion_matrix(test_generator.classes, np.where(preds[:,1]>0.
    ↪5,1,0), labels=[1, 0])
print(confusion)

```

```

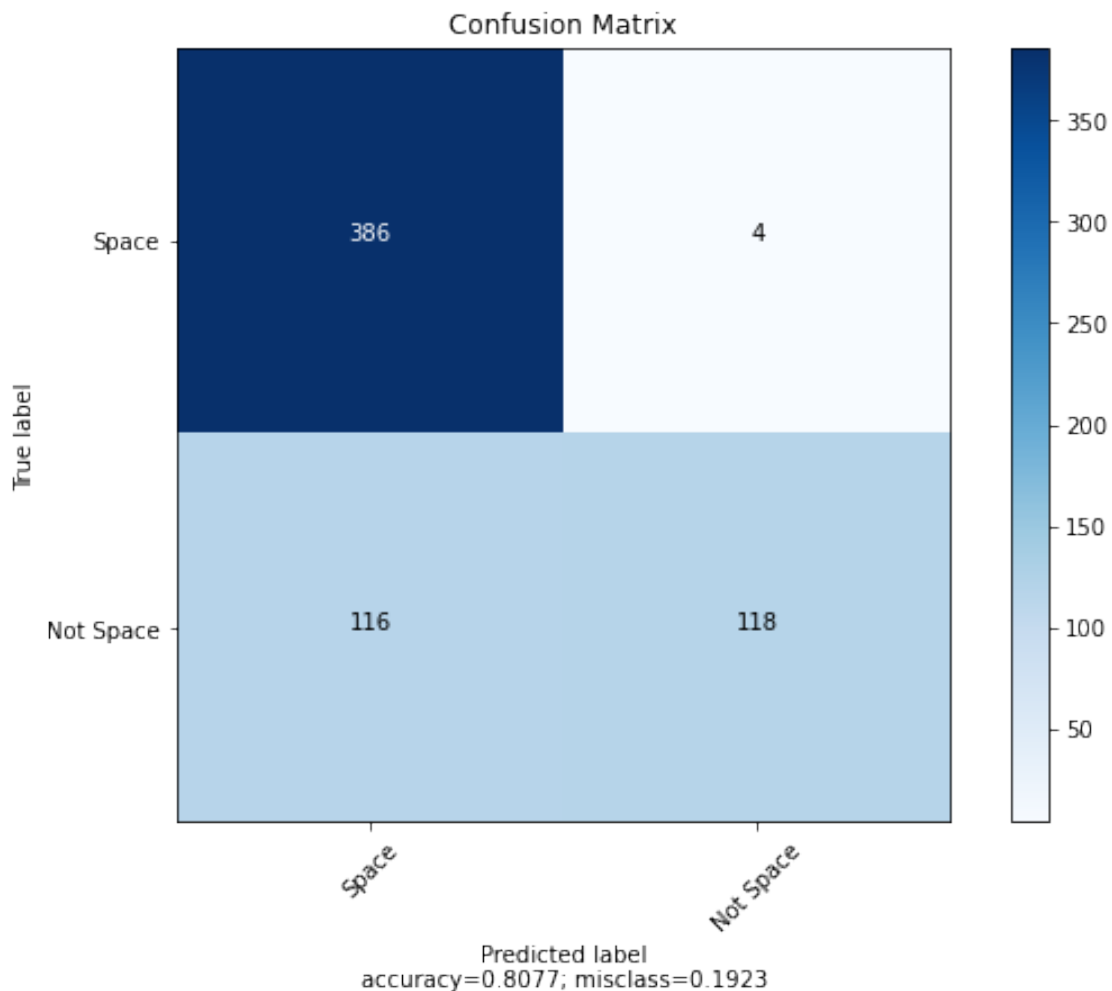
[[386   4]
 [116 118]]

```

```

[82]: plot_confusion_matrix(cm=confusion, target_names = ['Space', 'Not Space'],
    ↪title = 'Confusion Matrix',normalize=False)

```




```
[83]: from sklearn.metrics import classification_report

target_names = ['negative', 'positive']
print(classification_report(test_generator.classes, np.where(preds[:,1]>0.
↪5,1,0) , target_names=target_names))
```

	precision	recall	f1-score	support
negative	0.97	0.50	0.66	234
positive	0.77	0.99	0.87	390
accuracy			0.81	624
macro avg	0.87	0.75	0.76	624
weighted avg	0.84	0.81	0.79	624

```
[84]: fpr, tpr, _ = roc_curve(test_generator.classes, preds[:,1])
roc_auc = auc(fpr, tpr)
plt.figure()
lw = 2
plt.plot(fpr, tpr, color='darkorange',
lw=lw, label='ROC curve (area = %0.2f)' % roc_auc)
plt.plot([0, 1], [0, 1], color='navy', lw=lw, linestyle='--')
plt.xlim([-0.05, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic curve')
plt.legend(loc="lower right")
plt.show()
```



[]:

2 Xception

```
[121]: #data augmentation
train_gen = ImageDataGenerator(
    rescale=1/255.,
    featurewise_center=False, # set input mean to 0 over the dataset
    samplewise_center=False, # set each sample mean to 0
    featurewise_std_normalization=False, # divide inputs by std of the
    ↪ dataset
    samplewise_std_normalization=False, # divide each input by its std
    zca_whitening=False, # apply ZCA whitening
    rotation_range = 30, # randomly rotate images in the range (degrees, 0
    ↪ to 180)
    zoom_range = 0.4, # Randomly zoom image
    width_shift_range=0.1, # randomly shift images horizontally (fraction
    ↪ of total width)
    height_shift_range=0.1, # randomly shift images vertically (fraction
    ↪ of total height)
    horizontal_flip = True, # randomly flip images
    vertical_flip=False,
```

```

        validation_split=0.2) #validation size
test_gen=ImageDataGenerator(
    rescale=1/255.,
    featurewise_center=False, # set input mean to 0 over the dataset
    samplewise_center=False, # set each sample mean to 0
    featurewise_std_normalization=False, # divide inputs by std of the
    ↪dataset
    samplewise_std_normalization=False, # divide each input by its std
    zca_whitening=False, # apply ZCA whitening
    horizontal_flip = False, # randomly flip images
    vertical_flip=False) # randomly flip images

```

```

[122]: batch_size=16
nb_epochs=90
train_generator = train_gen.flow_from_directory(
    wd+"/train",
    batch_size=batch_size,
    target_size=(224, 224), #class_mode="binary"
    # class_mode='binary',
    subset='training'
)

validation_generator = train_gen.flow_from_directory(
    wd+"/train", # same directory as training data
    batch_size=batch_size,
    target_size=(224, 224),
    # class_mode='binary',
    subset='validation') # set as validation data
test_generator=test_gen.flow_from_directory(
    wd+"/test", # same directory as training data
    batch_size=1,
    target_size=(224, 224),
    shuffle=False)

```

Found 4173 images belonging to 2 classes.

Found 1043 images belonging to 2 classes.

Found 624 images belonging to 2 classes.

```

[123]: ##First, instantiate a base model with pre-trained weights.
base_model = tf.keras.applications.Xception(
    weights="imagenet", # Load weights pre-trained on ImageNet.
    input_shape=(224, 224, 3),
    include_top=False,
) # Do not include the ImageNet classifier at the top.

# Freeze the base_model
base_model.trainable = False

```



```

# Create new model on top
inputs = tf.keras.Input(shape=(224, 224, 3))
#x = tf.data_augmentation(inputs) # Apply random data augmentation

# Pre-trained Xception weights requires that input be normalized
# from (0, 255) to a range (-1., +1.), the normalization layer
# does the following, outputs = (inputs - mean) / sqrt(var)
norm_layer = tf.keras.layers.experimental.preprocessing.Normalization()
mean = np.array([0.5] * 3)
var = mean ** 2
# Scale inputs to [-1, +1]
x = norm_layer(inputs)
norm_layer.set_weights([mean, var])
# The base model contains batchnorm layers. We want to keep them in inference_
→mode
# when we unfreeze the base model for fine-tuning, so we make sure that the
# base_model is running in inference mode here.
x = base_model(x, training=False)
x = tf.keras.layers.GlobalAveragePooling2D()(x)
x = tf.keras.layers.Dropout(0.2)(x) # Regularize with dropout
outputs = tf.keras.layers.Dense(1)(x)
model_base = tf.keras.Model(inputs, outputs)

early_stopping_cb = tf.keras.callbacks.
→EarlyStopping(patience=10, restore_best_weights=True)
model=Sequential()
model.add(model_base)
model.add(Flatten())

model.add(Dense(128,activation='relu'))
model.add(Dense(64,activation='relu'))
model.add(Dense(2,activation='softmax'))

## Freezing the layers
#for layer in base_model.layers:
#    layer.trainable=False

model.compile(optimizer='adam', loss='categorical_crossentropy',_
→metrics=['accuracy'])

model_base.summary()
start=time.time()
history_xception = model.
→fit_generator(train_generator,epochs=90,validation_data=validation_generator,steps_per_epoch=
run_1=time.time()-start

```

Model: "model_4"

Layer (type)	Output Shape	Param #
input_11 (InputLayer)	[(None, 224, 224, 3)]	0
normalization_4 (Normalizati	(None, 224, 224, 3)	7
xception (Functional)	(None, 7, 7, 2048)	20861480
global_average_pooling2d_5 ((None, 2048)	0
dropout_8 (Dropout)	(None, 2048)	0
dense_34 (Dense)	(None, 1)	2049

Total params: 20,863,536

Trainable params: 2,049

Non-trainable params: 20,861,487

Epoch 1/90

260/260 [=====] - 74s 280ms/step - loss: 0.3883 - accuracy: 0.8114 - val_loss: 0.2270 - val_accuracy: 0.8993

Epoch 2/90

260/260 [=====] - 72s 278ms/step - loss: 0.2231 - accuracy: 0.9030 - val_loss: 0.2072 - val_accuracy: 0.9166

Epoch 3/90

260/260 [=====] - 72s 277ms/step - loss: 0.2061 - accuracy: 0.9141 - val_loss: 0.2020 - val_accuracy: 0.9166

Epoch 4/90

260/260 [=====] - 72s 279ms/step - loss: 0.2014 - accuracy: 0.9199 - val_loss: 0.1801 - val_accuracy: 0.9319

Epoch 5/90

260/260 [=====] - 72s 279ms/step - loss: 0.1951 - accuracy: 0.9223 - val_loss: 0.1834 - val_accuracy: 0.9204

Epoch 6/90

260/260 [=====] - 72s 277ms/step - loss: 0.1550 - accuracy: 0.9383 - val_loss: 0.1791 - val_accuracy: 0.9175

Epoch 7/90

260/260 [=====] - 72s 278ms/step - loss: 0.1779 - accuracy: 0.9296 - val_loss: 0.1696 - val_accuracy: 0.9281

Epoch 8/90

260/260 [=====] - 72s 278ms/step - loss: 0.1833 - accuracy: 0.9245 - val_loss: 0.1814 - val_accuracy: 0.9204

Epoch 9/90

260/260 [=====] - 73s 282ms/step - loss: 0.1806 - accuracy: 0.9264 - val_loss: 0.1994 - val_accuracy: 0.9156

Epoch 10/90

260/260 [=====] - 74s 284ms/step - loss: 0.1701 - accuracy: 0.9354 - val_loss: 0.1800 - val_accuracy: 0.9214
Epoch 11/90
260/260 [=====] - 72s 279ms/step - loss: 0.1630 - accuracy: 0.9407 - val_loss: 0.1774 - val_accuracy: 0.9300
Epoch 12/90
260/260 [=====] - 73s 279ms/step - loss: 0.1637 - accuracy: 0.9342 - val_loss: 0.1658 - val_accuracy: 0.9386
Epoch 13/90
260/260 [=====] - 72s 278ms/step - loss: 0.1683 - accuracy: 0.9394 - val_loss: 0.1547 - val_accuracy: 0.9415
Epoch 14/90
260/260 [=====] - 72s 278ms/step - loss: 0.1629 - accuracy: 0.9370 - val_loss: 0.1572 - val_accuracy: 0.9434
Epoch 15/90
260/260 [=====] - 72s 278ms/step - loss: 0.1587 - accuracy: 0.9352 - val_loss: 0.1608 - val_accuracy: 0.9367
Epoch 16/90
260/260 [=====] - 72s 278ms/step - loss: 0.1715 - accuracy: 0.9311 - val_loss: 0.1636 - val_accuracy: 0.9204
Epoch 17/90
260/260 [=====] - 73s 279ms/step - loss: 0.1804 - accuracy: 0.9261 - val_loss: 0.1833 - val_accuracy: 0.9243
Epoch 18/90
260/260 [=====] - 72s 279ms/step - loss: 0.1780 - accuracy: 0.9300 - val_loss: 0.1638 - val_accuracy: 0.9319
Epoch 19/90
260/260 [=====] - 72s 277ms/step - loss: 0.1598 - accuracy: 0.9418 - val_loss: 0.1752 - val_accuracy: 0.9348
Epoch 20/90
260/260 [=====] - 72s 277ms/step - loss: 0.1675 - accuracy: 0.9396 - val_loss: 0.1545 - val_accuracy: 0.9415
Epoch 21/90
260/260 [=====] - 72s 277ms/step - loss: 0.1610 - accuracy: 0.9427 - val_loss: 0.1649 - val_accuracy: 0.9358
Epoch 22/90
260/260 [=====] - 72s 278ms/step - loss: 0.1597 - accuracy: 0.9393 - val_loss: 0.1645 - val_accuracy: 0.9338
Epoch 23/90
260/260 [=====] - 72s 277ms/step - loss: 0.1448 - accuracy: 0.9432 - val_loss: 0.1347 - val_accuracy: 0.9521
Epoch 24/90
260/260 [=====] - 72s 277ms/step - loss: 0.1644 - accuracy: 0.9336 - val_loss: 0.1722 - val_accuracy: 0.9233
Epoch 25/90
260/260 [=====] - 72s 276ms/step - loss: 0.1801 - accuracy: 0.9311 - val_loss: 0.1673 - val_accuracy: 0.9291
Epoch 26/90

```

260/260 [=====] - 72s 277ms/step - loss: 0.1579 -
accuracy: 0.9398 - val_loss: 0.1403 - val_accuracy: 0.9473
Epoch 27/90
260/260 [=====] - 72s 277ms/step - loss: 0.1618 -
accuracy: 0.9419 - val_loss: 0.1684 - val_accuracy: 0.9358
Epoch 28/90
260/260 [=====] - 72s 277ms/step - loss: 0.1597 -
accuracy: 0.9343 - val_loss: 0.1443 - val_accuracy: 0.9415
Epoch 29/90
260/260 [=====] - 72s 277ms/step - loss: 0.1673 -
accuracy: 0.9296 - val_loss: 0.1438 - val_accuracy: 0.9406
Epoch 30/90
260/260 [=====] - 72s 277ms/step - loss: 0.1594 -
accuracy: 0.9352 - val_loss: 0.1320 - val_accuracy: 0.9463
Epoch 31/90
260/260 [=====] - 72s 278ms/step - loss: 0.1595 -
accuracy: 0.9363 - val_loss: 0.1655 - val_accuracy: 0.9348
Epoch 32/90
260/260 [=====] - 72s 277ms/step - loss: 0.1606 -
accuracy: 0.9398 - val_loss: 0.1575 - val_accuracy: 0.9396
Epoch 33/90
260/260 [=====] - 72s 277ms/step - loss: 0.1592 -
accuracy: 0.9370 - val_loss: 0.1595 - val_accuracy: 0.9348
Epoch 34/90
260/260 [=====] - 72s 277ms/step - loss: 0.1608 -
accuracy: 0.9434 - val_loss: 0.1371 - val_accuracy: 0.9501
Epoch 35/90
260/260 [=====] - 72s 277ms/step - loss: 0.1644 -
accuracy: 0.9386 - val_loss: 0.1613 - val_accuracy: 0.9386
Epoch 36/90
260/260 [=====] - 72s 277ms/step - loss: 0.1666 -
accuracy: 0.9291 - val_loss: 0.1588 - val_accuracy: 0.9386
Epoch 37/90
260/260 [=====] - 72s 279ms/step - loss: 0.1475 -
accuracy: 0.9425 - val_loss: 0.1609 - val_accuracy: 0.9367
Epoch 38/90
260/260 [=====] - 73s 280ms/step - loss: 0.1427 -
accuracy: 0.9425 - val_loss: 0.1511 - val_accuracy: 0.9348
Epoch 39/90
260/260 [=====] - 73s 279ms/step - loss: 0.1589 -
accuracy: 0.9378 - val_loss: 0.1521 - val_accuracy: 0.9425
Epoch 40/90
260/260 [=====] - 72s 277ms/step - loss: 0.1570 -
accuracy: 0.9420 - val_loss: 0.1627 - val_accuracy: 0.9329

```

```

[124]: # Unfreeze the base_model. Note that it keeps running in inference mode
# since we passed `training=False` when calling it. This means that

```

```

# the batchnorm layers will not update their batch statistics.
# This prevents the batchnorm layers from undoing all the training
# we've done so far.
base_model.trainable = True
model.summary()

model.compile(
    optimizer=tf.keras.optimizers.Adam(1e-5), # Low learning rate
    loss='categorical_crossentropy',
    metrics=['accuracy'],
)

epochs = 10

```

Model: "sequential_10"

Layer (type)	Output Shape	Param #
model_4 (Functional)	(None, 1)	20863536
flatten_10 (Flatten)	(None, 1)	0
dense_35 (Dense)	(None, 128)	256
dense_36 (Dense)	(None, 64)	8256
dense_37 (Dense)	(None, 2)	130

Total params: 20,872,178
 Trainable params: 20,817,643
 Non-trainable params: 54,535

```

[125]: start=time.time()
history_xception_tuning = model.
    ↪ fit_generator(train_generator, epochs=10, validation_data=validation_generator)
runtime=time.time()-start+run_1

```

Epoch 1/10
 261/261 [=====] - 90s 335ms/step - loss: 0.1605 - accuracy: 0.9443 - val_loss: 0.1123 - val_accuracy: 0.9501
 Epoch 2/10
 261/261 [=====] - 88s 335ms/step - loss: 0.1147 - accuracy: 0.9554 - val_loss: 0.1328 - val_accuracy: 0.9463
 Epoch 3/10
 261/261 [=====] - 88s 335ms/step - loss: 0.0929 - accuracy: 0.9653 - val_loss: 0.0914 - val_accuracy: 0.9664

```

Epoch 4/10
261/261 [=====] - 88s 337ms/step - loss: 0.0837 -
accuracy: 0.9693 - val_loss: 0.0811 - val_accuracy: 0.9703
Epoch 5/10
261/261 [=====] - 88s 338ms/step - loss: 0.0854 -
accuracy: 0.9664 - val_loss: 0.0745 - val_accuracy: 0.9693
Epoch 6/10
261/261 [=====] - 88s 337ms/step - loss: 0.0775 -
accuracy: 0.9691 - val_loss: 0.0560 - val_accuracy: 0.9789
Epoch 7/10
261/261 [=====] - 88s 338ms/step - loss: 0.0596 -
accuracy: 0.9765 - val_loss: 0.0603 - val_accuracy: 0.9770
Epoch 8/10
261/261 [=====] - 88s 337ms/step - loss: 0.0555 -
accuracy: 0.9781 - val_loss: 0.0725 - val_accuracy: 0.9732
Epoch 9/10
261/261 [=====] - 88s 336ms/step - loss: 0.0521 -
accuracy: 0.9813 - val_loss: 0.0701 - val_accuracy: 0.9770
Epoch 10/10
261/261 [=====] - 88s 337ms/step - loss: 0.0464 -
accuracy: 0.9844 - val_loss: 0.0714 - val_accuracy: 0.9760

```

```

[126]: score = model.evaluate(test_generator, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

```

```

Test score: 0.1737084835767746
Test accuracy: 0.932692289352417

```

```

[127]: #saving the model and histories
#tf.keras.models.save_model(filepath=cwd, model=model)
model.save(cwd+'/assets/model_xception')
np.save(cwd+'/assets/history_xception.npy', history_xception.history)
np.save(cwd+'/assets/history_xception_tuning.npy', history_xception_tuning.
    ↳history)
results['model'].append('Xception2')
results['total_param'].append(20863536)
results['time'].append(runtime)
results['score'].append(score[1])
np.save(cwd+'/assets/results.npy', results)

```

```

INFO:tensorflow:Assets written to:
/home/jovyan/work/fc_project/assets/model_xception/assets

```

```

[128]: fig, (ax1, ax2) = plt.subplots(1,2, figsize=(12,5))
ax1.plot(history_xception_tuning.history['loss'])
ax1.plot(history_xception_tuning.history['val_loss'])
ax1.plot(history_xception.history['loss'])


```

```

ax1.plot(history_xception.history['val_loss'])
ax1.set_title('loss')
ax1.set_ylabel('loss')
ax1.set_xlabel('epoch')
ax1.legend(['train tuned', 'test tuned', 'train', 'test'], loc='upper right')
ax2.plot(history_xception_tuning.history['accuracy'])
ax2.plot(history_xception_tuning.history['val_accuracy'])
ax2.plot(history_xception.history['accuracy'])
ax2.plot(history_xception.history['val_accuracy'])

ax2.set_title('accuracy')
ax2.set_ylabel('accuracy')
ax2.set_xlabel('epoch')
ax2.legend(['train tuned', 'test tuned', 'train', 'test'], loc='lower right')
fig.suptitle('loss and accuracy Xception model')
plt.show()

```



[]:

[]:

[]:

[]:

```

[129]: STEP_SIZE_TEST=test_generator.n//test_generator.batch_size
        test_generator.reset()
        preds = model.predict(test_generator,
                               steps=STEP_SIZE_TEST,

```

```
verbose=1)
```

624/624 [=====] - 6s 9ms/step

```
[130]: #scale the output of the model to [0,1]
#scaled_preds=(preds[:,1]-min(preds[:,1]))/(max(preds[:,1])-min(preds[:,1]))
pos = [i for i, j in zip(preds[:,1], test_generator.classes) if j == 1]
neg = [i for i, j in zip(preds[:,1], test_generator.classes) if j == 0]


with plt.xkcd():
    fig = plt.figure(figsize=(8, 4))

    sns.distplot(pos, hist = True, kde = False, color='g',
                  kde_kws = {'shade': True, 'linewidth': 3})

    sns.distplot(neg, hist = True, kde = False, color='r',
                  kde_kws = {'shade': True, 'linewidth': 3})

    plt.plot([0.5, 0.5], [0, 25], '-b')
    plt.annotate(
        'The probability threshold\npositive to the right\nnegative to the left',
        xy=(0.51, 15), arrowprops=dict(arrowstyle='->'), xytext=(0.6, 20))

plt.show()
```



```
[ ]:
```



```
[ ]:
```

```
[131]: confusion = confusion_matrix(test_generator.classes, np.where(preds[:,1]>0.
↳5,1,0), labels=[1, 0])
print(confusion)
```

```
[[369  21]
 [ 21 213]]
```

```
[132]: def plot_confusion_matrix(cm,
                                target_names,
                                title='Confusion matrix',
                                cmap=None,
                                normalize=True):

    """
    Given a scikit-learn confusion matrix (CM), make a nice plot.

    Arguments
    -----
    cm:          Confusion matrix from sklearn.metrics.confusion_matrix

    target_names: Given classification classes, such as [0, 1, 2]
                  The class names, for example, ['high', 'medium', 'low']

    title:       The text to display at the top of the matrix

    cmap:        The gradient of the values displayed from matplotlib.pyplot.cm
                  See http://matplotlib.org/examples/color/colormaps\_reference.
    ↳html
                  `plt.get_cmap('jet')` or `plt.cm.Blues`

    normalize:   If `False`, plot the raw numbers
                  If `True`, plot the proportions

    Usage
    ----
    plot_confusion_matrix(cm          = cm,                # Confusion_
    ↳matrix created by                                # `sklearn.
    ↳metrics.confusion_matrix`
                                normalize = True,          # Show proportions
                                target_names = y_labels_vals, # List of names_
    ↳of the classes
                                title      = best_estimator_name) # Title of graph

    Citation
```

```

-----
http://scikit-learn.org/stable/auto\_examples/model\_selection/
→ plot\_confusion\_matrix.html

"""
import matplotlib.pyplot as plt
import numpy as np
import itertools

accuracy = np.trace(cm) / float(np.sum(cm))
misclass = 1 - accuracy

if cmap is None:
    cmap = plt.get_cmap('Blues')
plt.figure(figsize=(8, 6))
plt.imshow(cm, interpolation='nearest', cmap=cmap)
plt.title(title)
plt.colorbar()

if target_names is not None:
    tick_marks = np.arange(len(target_names))
    plt.xticks(tick_marks, target_names, rotation=45)
    plt.yticks(tick_marks, target_names)


if normalize:
    cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]

thresh = cm.max() / 1.5 if normalize else cm.max() / 2
for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
    if normalize:
        plt.text(j, i, "{:0.4f}".format(cm[i, j]),
                  horizontalalignment="center",
                  color="white" if cm[i, j] > thresh else "black")
    else:
        plt.text(j, i, "{:,}".format(cm[i, j]),
                  horizontalalignment="center",
                  color="white" if cm[i, j] > thresh else "black")

plt.tight_layout()
plt.ylabel('True label')
plt.xlabel('Predicted label\naccuracy={:0.4f}; misclass={:0.4f}'.
→ format(accuracy, misclass))
plt.show()

```

```
[133]: plot_confusion_matrix(cm=confusion, target_names = ['Space', 'Not Space'],
    ↪title = 'Confusion Matrix',normalize=False)
```




```
[134]: from sklearn.metrics import classification_report

target_names = ['negative', 'positive']
print(classification_report(test_generator.classes, np.where(preds[:,1]>0.
    ↪5,1,0) , target_names=target_names))
```

	precision	recall	f1-score	support
negative	0.91	0.91	0.91	234
positive	0.95	0.95	0.95	390
accuracy			0.93	624
macro avg	0.93	0.93	0.93	624

weighted avg 0.93 0.93 0.93 624

```
[135]: fpr, tpr, _ = roc_curve(test_generator.classes, preds[:,1])
roc_auc = auc(fpr, tpr)
plt.figure()
lw = 2
plt.plot(fpr, tpr, color='darkorange',
lw=lw, label='ROC curve (area = %0.2f)' % roc_auc)
plt.plot([0, 1], [0, 1], color='navy', lw=lw, linestyle='--')
plt.xlim([-0.05, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic curve')
plt.legend(loc="lower right")
plt.show()
```



3 RESNET50

```
[186]: #data augmentation
train_gen = ImageDataGenerator(
    #    rescale=1/255,
    dtype = 'float32',
    #    featurewise_center=False, # set input mean to 0 over the dataset
    #    samplewise_center=False, # set each sample mean to 0
    #    featurewise_std_normalization=False, # divide inputs by std of the
    dataset
    #    samplewise_std_normalization=False, # divide each input by its std
    zca_whitening=False, # apply ZCA whitening
    rotation_range = 30, # randomly rotate images in the range (degrees, 0
    to 180)
    zoom_range = 0.4, # Randomly zoom image
    width_shift_range=0.1, # randomly shift images horizontally (fraction
    of total width)
    height_shift_range=0.1, # randomly shift images vertically (fraction
    of total height)
    horizontal_flip = True, # randomly flip images
    vertical_flip=False,
    preprocessing_function=tf.keras.applications.resnet50.preprocess_input,
    validation_split=0.2) #validation size
test_gen=ImageDataGenerator(
    #    rescale=1/255,
    dtype = 'float32',
    featurewise_center=False, # set input mean to 0 over the dataset
    samplewise_center=False, # set each sample mean to 0
    featurewise_std_normalization=False, # divide inputs by std of the
    dataset
    samplewise_std_normalization=False, # divide each input by its std
    zca_whitening=False, # apply ZCA whitening
    horizontal_flip = False, # randomly flip images
    preprocessing_function=tf.keras.applications.resnet50.preprocess_input,
    vertical_flip=False) # randomly flip images
```

```
[187]: batch_size=16
nb_epochs=90
train_generator = train_gen.flow_from_directory(
    wd+"/train",
    batch_size=batch_size,
    target_size=(224, 224), #class_mode="binary"
    #    class_mode='binary',
    subset='training'
)

validation_generator = train_gen.flow_from_directory(
```

```

wd+"/train", # same directory as training data
batch_size=batch_size,
target_size=(224, 224),
# class_mode='binary',
subset='validation') # set as validation data
test_generator=test_gen.flow_from_directory(
wd+"/test", # same directory as training data
batch_size=1,
target_size=(224, 224),
shuffle=False)

```

Found 4173 images belonging to 2 classes.
Found 1043 images belonging to 2 classes.
Found 624 images belonging to 2 classes.

```

[188]: ##First, instantiate a base model with pre-trained weights.
base_model = tf.keras.applications.ResNet50(
    weights="imagenet", # Load weights pre-trained on ImageNet.
    input_shape=(224, 224, 3),
    include_top=False,
) # Do not include the ImageNet classifier at the top.

# Freeze the base_model
base_model.trainable = False

# Create new model on top
inputs = tf.keras.Input(shape=(224, 224, 3))
#x = tf.data_augmentation(inputs) # Apply random data augmentation

x = base_model(inputs, training=False)
x = tf.keras.layers.GlobalAveragePooling2D()(x)
x = tf.keras.layers.Dropout(0.2)(x) # Regularize with dropout
outputs = tf.keras.layers.Dense(1)(x)
model_base = tf.keras.Model(inputs, outputs)

early_stopping_cb = tf.keras.callbacks.
    ↳EarlyStopping(patience=10,restore_best_weights=True)
model=Sequential()
model.add(model_base)
model.add(Flatten())

model.add(Dense(128,activation='relu'))
model.add(Dense(64,activation='relu'))
model.add(Dense(2,activation='softmax'))

## Freezing the layers

```

```

#for layer in base_model.layers:
#    layer.trainable=False

model.compile(optimizer='adam', loss='categorical_crossentropy',
↳metrics=['accuracy'])

model_base.summary()
start=time.time()
history_resnet50 = model.
↳fit_generator(train_generator,epochs=90,validation_data=validation_generator,steps_per_epoch=
run_1=time.time()-start

```

Model: "model_8"

Layer (type)	Output Shape	Param #
input_39 (InputLayer)	[(None, 224, 224, 3)]	0
resnet50 (Functional)	(None, 7, 7, 2048)	23587712
global_average_pooling2d_11	(None, 2048)	0
dropout_14 (Dropout)	(None, 2048)	0
dense_56 (Dense)	(None, 1)	2049

Total params: 23,589,761

Trainable params: 2,049

Non-trainable params: 23,587,712

Epoch 1/90

260/260 [=====] - 75s 282ms/step - loss: 0.4017 - accuracy: 0.7947 - val_loss: 0.1992 - val_accuracy: 0.9252

Epoch 2/90

260/260 [=====] - 72s 279ms/step - loss: 0.1732 - accuracy: 0.9302 - val_loss: 0.1591 - val_accuracy: 0.9348

Epoch 3/90

260/260 [=====] - 72s 278ms/step - loss: 0.1661 - accuracy: 0.9344 - val_loss: 0.1506 - val_accuracy: 0.9300

Epoch 4/90

260/260 [=====] - 72s 278ms/step - loss: 0.1392 - accuracy: 0.9486 - val_loss: 0.1261 - val_accuracy: 0.9511

Epoch 5/90

260/260 [=====] - 72s 277ms/step - loss: 0.1552 - accuracy: 0.9397 - val_loss: 0.1194 - val_accuracy: 0.9549

Epoch 6/90

260/260 [=====] - 72s 277ms/step - loss: 0.1338 -

accuracy: 0.9452 - val_loss: 0.1170 - val_accuracy: 0.9569
 Epoch 7/90
 260/260 [=====] - 72s 278ms/step - loss: 0.1255 -
 accuracy: 0.9525 - val_loss: 0.1055 - val_accuracy: 0.9578
 Epoch 8/90
 260/260 [=====] - 72s 278ms/step - loss: 0.1263 -
 accuracy: 0.9524 - val_loss: 0.1044 - val_accuracy: 0.9626
 Epoch 9/90
 260/260 [=====] - 72s 278ms/step - loss: 0.1212 -
 accuracy: 0.9532 - val_loss: 0.1061 - val_accuracy: 0.9530
 Epoch 10/90
 260/260 [=====] - 72s 277ms/step - loss: 0.1177 -
 accuracy: 0.9603 - val_loss: 0.1085 - val_accuracy: 0.9607
 Epoch 11/90
 260/260 [=====] - 72s 277ms/step - loss: 0.1375 -
 accuracy: 0.9427 - val_loss: 0.1081 - val_accuracy: 0.9626
 Epoch 12/90
 260/260 [=====] - 72s 277ms/step - loss: 0.1136 -
 accuracy: 0.9562 - val_loss: 0.0913 - val_accuracy: 0.9712
 Epoch 13/90
 260/260 [=====] - 72s 277ms/step - loss: 0.1153 -
 accuracy: 0.9568 - val_loss: 0.1119 - val_accuracy: 0.9540
 Epoch 14/90
 260/260 [=====] - 72s 277ms/step - loss: 0.1083 -
 accuracy: 0.9545 - val_loss: 0.1041 - val_accuracy: 0.9540
 Epoch 15/90
 260/260 [=====] - 73s 281ms/step - loss: 0.1169 -
 accuracy: 0.9524 - val_loss: 0.0971 - val_accuracy: 0.9616
 Epoch 16/90
 260/260 [=====] - 74s 286ms/step - loss: 0.1179 -
 accuracy: 0.9560 - val_loss: 0.1026 - val_accuracy: 0.9578
 Epoch 17/90
 260/260 [=====] - 73s 280ms/step - loss: 0.1209 -
 accuracy: 0.9541 - val_loss: 0.1350 - val_accuracy: 0.9492
 Epoch 18/90
 260/260 [=====] - 73s 282ms/step - loss: 0.1088 -
 accuracy: 0.9575 - val_loss: 0.1516 - val_accuracy: 0.9367
 Epoch 19/90
 260/260 [=====] - 73s 282ms/step - loss: 0.1153 -
 accuracy: 0.9591 - val_loss: 0.1133 - val_accuracy: 0.9521
 Epoch 20/90
 260/260 [=====] - 73s 279ms/step - loss: 0.1120 -
 accuracy: 0.9548 - val_loss: 0.1071 - val_accuracy: 0.9607
 Epoch 21/90
 260/260 [=====] - 72s 276ms/step - loss: 0.1207 -
 accuracy: 0.9535 - val_loss: 0.1096 - val_accuracy: 0.9559
 Epoch 22/90
 260/260 [=====] - 72s 277ms/step - loss: 0.1090 -

accuracy: 0.9548 - val_loss: 0.0897 - val_accuracy: 0.9674
 Epoch 23/90
 260/260 [=====] - 71s 274ms/step - loss: 0.1131 -
 accuracy: 0.9553 - val_loss: 0.1333 - val_accuracy: 0.9463
 Epoch 24/90
 260/260 [=====] - 72s 278ms/step - loss: 0.1126 -
 accuracy: 0.9558 - val_loss: 0.0963 - val_accuracy: 0.9569
 Epoch 25/90
 260/260 [=====] - 72s 277ms/step - loss: 0.1278 -
 accuracy: 0.9467 - val_loss: 0.1056 - val_accuracy: 0.9588
 Epoch 26/90
 260/260 [=====] - 73s 279ms/step - loss: 0.1069 -
 accuracy: 0.9557 - val_loss: 0.0894 - val_accuracy: 0.9645
 Epoch 27/90
 260/260 [=====] - 73s 282ms/step - loss: 0.1020 -
 accuracy: 0.9625 - val_loss: 0.0924 - val_accuracy: 0.9655
 Epoch 28/90
 260/260 [=====] - 72s 276ms/step - loss: 0.1020 -
 accuracy: 0.9627 - val_loss: 0.0959 - val_accuracy: 0.9626
 Epoch 29/90
 260/260 [=====] - 72s 276ms/step - loss: 0.1032 -
 accuracy: 0.9657 - val_loss: 0.0810 - val_accuracy: 0.9732
 Epoch 30/90
 260/260 [=====] - 71s 273ms/step - loss: 0.1011 -
 accuracy: 0.9595 - val_loss: 0.0853 - val_accuracy: 0.9684
 Epoch 31/90
 260/260 [=====] - 71s 273ms/step - loss: 0.1289 -
 accuracy: 0.9537 - val_loss: 0.1050 - val_accuracy: 0.9655
 Epoch 32/90
 260/260 [=====] - 73s 280ms/step - loss: 0.1137 -
 accuracy: 0.9569 - val_loss: 0.0862 - val_accuracy: 0.9655
 Epoch 33/90
 260/260 [=====] - 71s 275ms/step - loss: 0.0990 -
 accuracy: 0.9623 - val_loss: 0.1008 - val_accuracy: 0.9664
 Epoch 34/90
 260/260 [=====] - 71s 272ms/step - loss: 0.1069 -
 accuracy: 0.9598 - val_loss: 0.0972 - val_accuracy: 0.9626
 Epoch 35/90
 260/260 [=====] - 71s 273ms/step - loss: 0.1059 -
 accuracy: 0.9596 - val_loss: 0.1137 - val_accuracy: 0.9607
 Epoch 36/90
 260/260 [=====] - 72s 275ms/step - loss: 0.1183 -
 accuracy: 0.9567 - val_loss: 0.0892 - val_accuracy: 0.9655
 Epoch 37/90
 260/260 [=====] - 71s 273ms/step - loss: 0.1100 -
 accuracy: 0.9548 - val_loss: 0.0984 - val_accuracy: 0.9626
 Epoch 38/90
 260/260 [=====] - 71s 272ms/step - loss: 0.1005 -

```

accuracy: 0.9624 - val_loss: 0.1070 - val_accuracy: 0.9588
Epoch 39/90
260/260 [=====] - 71s 273ms/step - loss: 0.1125 -
accuracy: 0.9593 - val_loss: 0.0789 - val_accuracy: 0.9693
Epoch 40/90
260/260 [=====] - 71s 273ms/step - loss: 0.1026 -
accuracy: 0.9590 - val_loss: 0.1092 - val_accuracy: 0.9569
Epoch 41/90
260/260 [=====] - 73s 281ms/step - loss: 0.1042 -
accuracy: 0.9648 - val_loss: 0.1241 - val_accuracy: 0.9453
Epoch 42/90
260/260 [=====] - 72s 276ms/step - loss: 0.1018 -
accuracy: 0.9645 - val_loss: 0.0802 - val_accuracy: 0.9703
Epoch 43/90
260/260 [=====] - 71s 274ms/step - loss: 0.1135 -
accuracy: 0.9566 - val_loss: 0.0964 - val_accuracy: 0.9655
Epoch 44/90
260/260 [=====] - 71s 273ms/step - loss: 0.1082 -
accuracy: 0.9627 - val_loss: 0.0842 - val_accuracy: 0.9703
Epoch 45/90
260/260 [=====] - 71s 275ms/step - loss: 0.1004 -
accuracy: 0.9635 - val_loss: 0.0860 - val_accuracy: 0.9664
Epoch 46/90
260/260 [=====] - 71s 274ms/step - loss: 0.1029 -
accuracy: 0.9586 - val_loss: 0.1046 - val_accuracy: 0.9607
Epoch 47/90
260/260 [=====] - 72s 275ms/step - loss: 0.1147 -
accuracy: 0.9573 - val_loss: 0.0965 - val_accuracy: 0.9626
Epoch 48/90
260/260 [=====] - 71s 273ms/step - loss: 0.0997 -
accuracy: 0.9633 - val_loss: 0.1320 - val_accuracy: 0.9434
Epoch 49/90
260/260 [=====] - 71s 273ms/step - loss: 0.0923 -
accuracy: 0.9602 - val_loss: 0.0954 - val_accuracy: 0.9626

```

```

[189]: #saving the model and histories
#tf.keras.models.save_model(filepath=cwd, model=model)
model.save(cwd+'/assets/model_resnet50')
np.save(cwd+'/assets/history_resnet50.npy', history_resnet50.history)
#np.save(cwd+'/assets/history_xception_tuning.npy', history_xception_tuning.
↪history)

```

```

INFO:tensorflow:Assets written to:
/home/jovyan/work/fc_project/assets/model_resnet50/assets

```

```

[190]: # Unfreeze the base_model. Note that it keeps running in inference mode
# since we passed `training=False` when calling it. This means that

```

```

# the batchnorm layers will not update their batch statistics.
# This prevents the batchnorm layers from undoing all the training
# we've done so far.
base_model.trainable = True
model.summary()

model.compile(
    optimizer=tf.keras.optimizers.Adam(1e-5), # Low learning rate
    loss='categorical_crossentropy',
    metrics=['accuracy'],
)

epochs = 10

```

Model: "sequential_16"

Layer (type)	Output Shape	Param #
model_8 (Functional)	(None, 1)	23589761
flatten_16 (Flatten)	(None, 1)	0
dense_57 (Dense)	(None, 128)	256
dense_58 (Dense)	(None, 64)	8256
dense_59 (Dense)	(None, 2)	130

Total params: 23,598,403
 Trainable params: 23,545,283
 Non-trainable params: 53,120

```

[191]: start=time.time()
history_resnet50_tuning = model.
    ↪ fit_generator(train_generator, epochs=10, validation_data=validation_generator)
runtime=time.time()-start+run_1

```

Epoch 1/10
 261/261 [=====] - 77s 279ms/step - loss: 0.1480 - accuracy: 0.9447 - val_loss: 0.1170 - val_accuracy: 0.9569
 Epoch 2/10
 261/261 [=====] - 72s 276ms/step - loss: 0.1225 - accuracy: 0.9509 - val_loss: 0.0821 - val_accuracy: 0.9693
 Epoch 3/10
 261/261 [=====] - 72s 276ms/step - loss: 0.0821 - accuracy: 0.9696 - val_loss: 0.0694 - val_accuracy: 0.9732

```

Epoch 4/10
261/261 [=====] - 72s 277ms/step - loss: 0.0779 -
accuracy: 0.9692 - val_loss: 0.0741 - val_accuracy: 0.9712
Epoch 5/10
261/261 [=====] - 73s 277ms/step - loss: 0.0531 -
accuracy: 0.9805 - val_loss: 0.1266 - val_accuracy: 0.9540
Epoch 6/10
261/261 [=====] - 73s 278ms/step - loss: 0.0736 -
accuracy: 0.9712 - val_loss: 0.0534 - val_accuracy: 0.9789
Epoch 7/10
261/261 [=====] - 73s 277ms/step - loss: 0.0492 -
accuracy: 0.9839 - val_loss: 0.1070 - val_accuracy: 0.9540
Epoch 8/10
261/261 [=====] - 73s 279ms/step - loss: 0.0347 -
accuracy: 0.9866 - val_loss: 0.0537 - val_accuracy: 0.9808
Epoch 9/10
261/261 [=====] - 73s 278ms/step - loss: 0.0617 -
accuracy: 0.9789 - val_loss: 0.0469 - val_accuracy: 0.9818
Epoch 10/10
261/261 [=====] - 73s 278ms/step - loss: 0.0470 -
accuracy: 0.9844 - val_loss: 0.0750 - val_accuracy: 0.9760

```

```

[192]: #saving the model and histories
#tf.keras.models.save_model(filepath=cwd, model=model)
model.save(cwd+'/assets/model_resnet50_tuning')
np.save(cwd+'/assets/history_resnet50_tuning.npy', history_resnet50_tuning.
    ↳history)
#np.save(cwd+'/assets/history_xception_tuning.npy', history_xception_tuning.
    ↳history)

```

```

INFO:tensorflow:Assets written to:
/home/jovyan/work/fc_project/assets/model_resnet50_tuning/assets

```

```

[201]: model=tf.keras.models.load_model(cwd+'/assets/model_resnet50_tuning')
history_resnet50_tuning=np.load(cwd+'/assets/history_resnet50_tuning.npy',
    ↳allow_pickle=True)
history_resnet50=np.load(cwd+'/assets/history_resnet50.npy', allow_pickle=True)

```

```

[194]:

```

```

[202]: score = model.evaluate(test_generator, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

```

```

Test score: 0.2513222396373749
Test accuracy: 0.9070512652397156

```

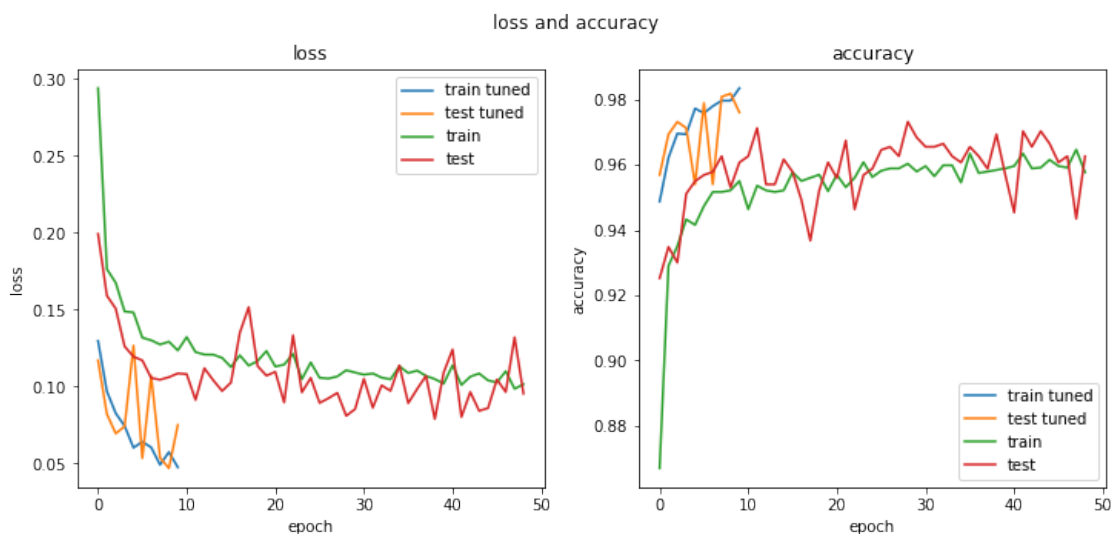
```
[203]: score = model_resnet50.evaluate(test_generator, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])
```

Test score: 0.21405178308486938
Test accuracy: 0.9086538553237915

```
[204]: history_resnet50_tuning=history_resnet50_tuning.tolist()
history_resnet50=history_resnet50.tolist()
```

```
[205]: fig, (ax1, ax2) = plt.subplots(1,2, figsize=(12,5))
ax1.plot(history_resnet50_tuning['loss'])
ax1.plot(history_resnet50_tuning['val_loss'])
ax1.plot(history_resnet50['loss'])
ax1.plot(history_resnet50['val_loss'])
ax1.set_title('loss')
ax1.set_ylabel('loss')
ax1.set_xlabel('epoch')
ax1.legend(['train tuned', 'test tuned', 'train', 'test'], loc='upper right')
ax2.plot(history_resnet50_tuning['accuracy'])
ax2.plot(history_resnet50_tuning['val_accuracy'])
ax2.plot(history_resnet50['accuracy'])
ax2.plot(history_resnet50['val_accuracy'])

ax2.set_title('accuracy')
ax2.set_ylabel('accuracy')
ax2.set_xlabel('epoch')
ax2.legend(['train tuned', 'test tuned', 'train', 'test'], loc='lower right')
fig.suptitle('loss and accuracy')
plt.show()
```



```
[206]: results['model'].append('resnet50')
results['total_param'].append(23598403)
results['time'].append(runtime)
results['score'].append(score[1])
np.save(cwd+'/assets/results.npy', results)
```

```
[207]: STEP_SIZE_TEST=test_generator.n//test_generator.batch_size
test_generator.reset()
preds = model.predict(test_generator,
steps=STEP_SIZE_TEST,
verbose=1)
```

624/624 [=====] - 6s 9ms/step

```
[208]: #scale the output of the model to [0,1]
#scaled_preds=(preds[:,1]-min(preds[:,1]))/(max(preds[:,1])-min(preds[:,1]))
pos = [i for i, j in zip(preds[:,1], test_generator.classes) if j == 1]
neg = [i for i, j in zip(preds[:,1], test_generator.classes) if j == 0]


with plt.xkcd():
    fig = plt.figure(figsize=(8, 4))

    sns.distplot(pos, hist = True, kde = False, color='g',
                  kde_kws = {'shade': True, 'linewidth': 3})

    sns.distplot(neg, hist = True, kde = False, color='r',
                  kde_kws = {'shade': True, 'linewidth': 3})

    plt.plot([0.5, 0.5], [0, 25], '-b')
    plt.annotate(
        'The probability threshold\npositive to the right\nnegative to the
↪left',
        xy=(0.51, 15), arrowprops=dict(arrowstyle='->'), xytext=(0.6, 20))


plt.show()
```



```
[209]: confusion = confusion_matrix(test_generator.classes, np.where(preds[:,1]>0.
    ↪5,1,0), labels=[1, 0])
print(confusion)
```

```
[[356  34]
 [ 24 210]]
```

```
[210]: plot_confusion_matrix(cm=confusion, target_names = ['Space', 'Not Space'],
    ↪title = 'Confusion Matrix',normalize=False)
```




```
[211]: target_names = ['negative', 'positive']
print(classification_report(test_generator.classes, np.where(preds[:,1]>0.
↪5,1,0) , target_names=target_names))
```

	precision	recall	f1-score	support
negative	0.86	0.90	0.88	234
positive	0.94	0.91	0.92	390
accuracy			0.91	624
macro avg	0.90	0.91	0.90	624
weighted avg	0.91	0.91	0.91	624

```
[212]: fpr, tpr, _ = roc_curve(test_generator.classes, preds[:,1])
roc_auc = auc(fpr, tpr)
```



```
plt.figure()
lw = 2
plt.plot(fpr, tpr, color='darkorange',
lw=lw, label='ROC curve (area = %0.2f)' % roc_auc)
plt.plot([0, 1], [0, 1], color='navy', lw=lw, linestyle='--')
plt.xlim([-0.05, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic curve')
plt.legend(loc="lower right")
plt.show()
```



4 VGG16

```
[213]: #data augmentation
train_gen = ImageDataGenerator(
    dtype = 'float32',
    # featurewise_center=False, # set input mean to 0 over the dataset
    # samplewise_center=False, # set each sample mean to 0
    # featurewise_std_normalization=False, # divide inputs by std of the
    ↪ dataset
```

```

# samplewise_std_normalization=False, # divide each input by its std
zca_whitening=False, # apply ZCA whitening
rotation_range = 30, # randomly rotate images in the range (degrees, 0
→to 180)
zoom_range = 0.4, # Randomly zoom image
width_shift_range=0.1, # randomly shift images horizontally (fraction
→of total width)
height_shift_range=0.1, # randomly shift images vertically (fraction
→of total height)
horizontal_flip = True, # randomly flip images
vertical_flip=False,
preprocessing_function=tf.keras.applications.vgg16.preprocess_input,
validation_split=0.2) #validation size
test_gen=ImageDataGenerator(
    dtype = 'float32',
    featurewise_center=False, # set input mean to 0 over the dataset
    samplewise_center=False, # set each sample mean to 0
    featurewise_std_normalization=False, # divide inputs by std of the
→dataset
    samplewise_std_normalization=False, # divide each input by its std
    zca_whitening=False, # apply ZCA whitening
    horizontal_flip = False, # randomly flip images
    preprocessing_function=tf.keras.applications.vgg16.preprocess_input,
    vertical_flip=False) # randomly flip images

```

```

[214]: batch_size=64
nb_epochs=90
train_generator = train_gen.flow_from_directory(
    wd+"/train",
    batch_size=batch_size,
    target_size=(224, 224), #class_mode="binary"
    # class_mode='binary',
    subset='training'
)

validation_generator = train_gen.flow_from_directory(
    wd+"/train", # same directory as training data
    batch_size=batch_size,
    target_size=(224, 224),
    # class_mode='binary',
    subset='validation') # set as validation data
test_generator=test_gen.flow_from_directory(
    wd+"/test", # same directory as training data
    batch_size=1,
    target_size=(224, 224),
    shuffle=False)

```

Found 4173 images belonging to 2 classes.
 Found 1043 images belonging to 2 classes.
 Found 624 images belonging to 2 classes.

```
[215]: inputs = tf.keras.Input(shape=(224,224,3))
##First, instantiate a base model with pre-trained weights.
base_model = tf.keras.applications.VGG16(
    weights="imagenet", # Load weights pre-trained on ImageNet.
    input_shape=(224,224,3),
    include_top=False,
    pooling='avg',
    input_tensor=inputs
) # Do not include the ImageNet classifier at the top.

# Freeze the base_model
base_model.trainable = False

# Create new model on top

early_stopping_cb = tf.keras.callbacks.
↳EarlyStopping(patience=10,restore_best_weights=True)
model=Sequential()
model.add(base_model)
model.add(Flatten())
model.add(BatchNormalization())
model.add(Dense(128,activation='relu'))
model.add(Dropout(0.5))
model.add(BatchNormalization())
model.add(Dense(64,activation='relu'))
model.add(Dropout(0.5))
model.add(BatchNormalization())
model.add(Dense(2,activation='softmax'))

model.summary()
model.compile(optimizer='adam', loss='categorical_crossentropy',
↳metrics=['accuracy'])
start=time.time()
history_vgg16 = model.
↳fit_generator(train_generator,epochs=90,validation_data=validation_generator,callbacks=[ear
runtime=time.time()-start
```

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/vgg16/vgg16_weights_tf_dim_ordering_tf_kernels_notop.h5
 58892288/58889256 [=====] - 138s 2us/step
 Model: "sequential_17"

Layer (type)	Output Shape	Param #
--------------	--------------	---------

```

=====
vgg16 (Functional)          (None, 512)          14714688
-----
flatten_17 (Flatten)        (None, 512)           0
-----
batch_normalization_33 (Batc (None, 512)          2048
-----
dense_60 (Dense)            (None, 128)          65664
-----
dropout_15 (Dropout)        (None, 128)           0
-----
batch_normalization_34 (Batc (None, 128)          512
-----
dense_61 (Dense)            (None, 64)           8256
-----
dropout_16 (Dropout)        (None, 64)           0
-----
batch_normalization_35 (Batc (None, 64)           256
-----
dense_62 (Dense)            (None, 2)            130
=====

```

Total params: 14,791,554

Trainable params: 75,458

Non-trainable params: 14,716,096

Epoch 1/90

2021-08-15 16:30:57.639008: W

tensorflow/core/common_runtime/bfc_allocator.cc:248] Allocator (GPU_0_bfc) ran out of memory trying to allocate 3.46GiB with freed_by_count=0. The caller indicates that this is not a failure, but may mean that there could be performance gains if more memory were available.

66/66 [=====] - 106s 1s/step - loss: 0.7699 - accuracy: 0.6566 - val_loss: 1.1946 - val_accuracy: 0.6261

Epoch 2/90

66/66 [=====] - 73s 1s/step - loss: 0.3752 - accuracy: 0.8584 - val_loss: 0.3569 - val_accuracy: 0.8792

Epoch 3/90

66/66 [=====] - 74s 1s/step - loss: 0.2759 - accuracy: 0.8960 - val_loss: 0.2283 - val_accuracy: 0.9271

Epoch 4/90

66/66 [=====] - 74s 1s/step - loss: 0.2371 - accuracy: 0.9103 - val_loss: 0.1751 - val_accuracy: 0.9348

Epoch 5/90

66/66 [=====] - 73s 1s/step - loss: 0.2212 - accuracy: 0.9171 - val_loss: 0.1488 - val_accuracy: 0.9473

Epoch 6/90

66/66 [=====] - 72s 1s/step - loss: 0.2149 - accuracy:

0.9177 - val_loss: 0.1332 - val_accuracy: 0.9511
 Epoch 7/90
 66/66 [=====] - 72s 1s/step - loss: 0.1989 - accuracy:
 0.9281 - val_loss: 0.1420 - val_accuracy: 0.9425
 Epoch 8/90
 66/66 [=====] - 72s 1s/step - loss: 0.1838 - accuracy:
 0.9296 - val_loss: 0.1315 - val_accuracy: 0.9511
 Epoch 9/90
 66/66 [=====] - 72s 1s/step - loss: 0.1782 - accuracy:
 0.9325 - val_loss: 0.1467 - val_accuracy: 0.9396
 Epoch 10/90
 66/66 [=====] - 72s 1s/step - loss: 0.1616 - accuracy:
 0.9404 - val_loss: 0.1388 - val_accuracy: 0.9530
 Epoch 11/90
 66/66 [=====] - 72s 1s/step - loss: 0.1633 - accuracy:
 0.9388 - val_loss: 0.1388 - val_accuracy: 0.9453
 Epoch 12/90
 66/66 [=====] - 73s 1s/step - loss: 0.1535 - accuracy:
 0.9431 - val_loss: 0.1284 - val_accuracy: 0.9530
 Epoch 13/90
 66/66 [=====] - 72s 1s/step - loss: 0.1680 - accuracy:
 0.9376 - val_loss: 0.1217 - val_accuracy: 0.9540
 Epoch 14/90
 66/66 [=====] - 72s 1s/step - loss: 0.1665 - accuracy:
 0.9398 - val_loss: 0.1325 - val_accuracy: 0.9549
 Epoch 15/90
 66/66 [=====] - 73s 1s/step - loss: 0.1621 - accuracy:
 0.9418 - val_loss: 0.1224 - val_accuracy: 0.9549
 Epoch 16/90
 66/66 [=====] - 73s 1s/step - loss: 0.1530 - accuracy:
 0.9406 - val_loss: 0.1328 - val_accuracy: 0.9559
 Epoch 17/90
 66/66 [=====] - 72s 1s/step - loss: 0.1418 - accuracy:
 0.9470 - val_loss: 0.1255 - val_accuracy: 0.9511
 Epoch 18/90
 66/66 [=====] - 72s 1s/step - loss: 0.1408 - accuracy:
 0.9446 - val_loss: 0.1232 - val_accuracy: 0.9540
 Epoch 19/90
 66/66 [=====] - 72s 1s/step - loss: 0.1537 - accuracy:
 0.9407 - val_loss: 0.1174 - val_accuracy: 0.9521
 Epoch 20/90
 66/66 [=====] - 72s 1s/step - loss: 0.1314 - accuracy:
 0.9536 - val_loss: 0.1148 - val_accuracy: 0.9549
 Epoch 21/90
 66/66 [=====] - 72s 1s/step - loss: 0.1424 - accuracy:
 0.9455 - val_loss: 0.1128 - val_accuracy: 0.9569
 Epoch 22/90
 66/66 [=====] - 72s 1s/step - loss: 0.1463 - accuracy:

0.9469 - val_loss: 0.1278 - val_accuracy: 0.9492
 Epoch 23/90
 66/66 [=====] - 72s 1s/step - loss: 0.1477 - accuracy:
 0.9492 - val_loss: 0.1137 - val_accuracy: 0.9530
 Epoch 24/90
 66/66 [=====] - 72s 1s/step - loss: 0.1337 - accuracy:
 0.9490 - val_loss: 0.1028 - val_accuracy: 0.9626
 Epoch 25/90
 66/66 [=====] - 72s 1s/step - loss: 0.1551 - accuracy:
 0.9425 - val_loss: 0.1168 - val_accuracy: 0.9626
 Epoch 26/90
 66/66 [=====] - 72s 1s/step - loss: 0.1334 - accuracy:
 0.9515 - val_loss: 0.1012 - val_accuracy: 0.9559
 Epoch 27/90
 66/66 [=====] - 71s 1s/step - loss: 0.1312 - accuracy:
 0.9503 - val_loss: 0.1065 - val_accuracy: 0.9597
 Epoch 28/90
 66/66 [=====] - 71s 1s/step - loss: 0.1147 - accuracy:
 0.9585 - val_loss: 0.1158 - val_accuracy: 0.9530
 Epoch 29/90
 66/66 [=====] - 72s 1s/step - loss: 0.1310 - accuracy:
 0.9509 - val_loss: 0.1078 - val_accuracy: 0.9578
 Epoch 30/90
 66/66 [=====] - 71s 1s/step - loss: 0.1199 - accuracy:
 0.9517 - val_loss: 0.1149 - val_accuracy: 0.9588
 Epoch 31/90
 66/66 [=====] - 72s 1s/step - loss: 0.1134 - accuracy:
 0.9566 - val_loss: 0.0987 - val_accuracy: 0.9655
 Epoch 32/90
 66/66 [=====] - 72s 1s/step - loss: 0.1218 - accuracy:
 0.9517 - val_loss: 0.1011 - val_accuracy: 0.9569
 Epoch 33/90
 66/66 [=====] - 72s 1s/step - loss: 0.1179 - accuracy:
 0.9592 - val_loss: 0.1017 - val_accuracy: 0.9588
 Epoch 34/90
 66/66 [=====] - 71s 1s/step - loss: 0.1263 - accuracy:
 0.9532 - val_loss: 0.1006 - val_accuracy: 0.9636
 Epoch 35/90
 66/66 [=====] - 72s 1s/step - loss: 0.1247 - accuracy:
 0.9517 - val_loss: 0.0957 - val_accuracy: 0.9655
 Epoch 36/90
 66/66 [=====] - 72s 1s/step - loss: 0.1234 - accuracy:
 0.9533 - val_loss: 0.1059 - val_accuracy: 0.9569
 Epoch 37/90
 66/66 [=====] - 72s 1s/step - loss: 0.1159 - accuracy:
 0.9572 - val_loss: 0.1182 - val_accuracy: 0.9540
 Epoch 38/90
 66/66 [=====] - 72s 1s/step - loss: 0.1087 - accuracy:

```

0.9582 - val_loss: 0.1030 - val_accuracy: 0.9597
Epoch 39/90
66/66 [=====] - 72s 1s/step - loss: 0.1269 - accuracy:
0.9551 - val_loss: 0.1209 - val_accuracy: 0.9521
Epoch 40/90
66/66 [=====] - 72s 1s/step - loss: 0.1278 - accuracy:
0.9574 - val_loss: 0.1079 - val_accuracy: 0.9588
Epoch 41/90
66/66 [=====] - 72s 1s/step - loss: 0.1174 - accuracy:
0.9548 - val_loss: 0.1202 - val_accuracy: 0.9492
Epoch 42/90
66/66 [=====] - 72s 1s/step - loss: 0.1365 - accuracy:
0.9443 - val_loss: 0.0997 - val_accuracy: 0.9607
Epoch 43/90
66/66 [=====] - 72s 1s/step - loss: 0.1311 - accuracy:
0.9488 - val_loss: 0.1025 - val_accuracy: 0.9664
Epoch 44/90
66/66 [=====] - 72s 1s/step - loss: 0.1114 - accuracy:
0.9582 - val_loss: 0.1092 - val_accuracy: 0.9549
Epoch 45/90
66/66 [=====] - 72s 1s/step - loss: 0.1081 - accuracy:
0.9584 - val_loss: 0.1129 - val_accuracy: 0.9645

```

```

[216]: #saving the model and histories
#tf.keras.models.save_model(filepath=cwd, model=model)
model.save(cwd+'/assets/model_vgg16')
np.save(cwd+'/assets/history_vgg16.npy', history_vgg16.history)
#np.save(cwd+'/assets/history_xception_tuning.npy', history_xception_tuning.
↪history)

```

```

INFO:tensorflow:Assets written to:
/home/jovyan/work/fc_project/assets/model_vgg16/assets

```

```

[217]: #tuning
# Unfreeze the base_model. Note that it keeps running in inference mode
# since we passed `training=False` when calling it. This means that
# the batchnorm layers will not update their batch statistics.
# This prevents the batchnorm layers from undoing all the training
# we've done so far.
base_model.trainable = True
model.summary()

model.compile(
    optimizer=tf.keras.optimizers.Adam(1e-5), # Low learning rate
    loss='categorical_crossentropy',
    metrics=['accuracy'],
)

```

```
epochs = 10
```

Model: "sequential_17"

Layer (type)	Output Shape	Param #
vgg16 (Functional)	(None, 512)	14714688
flatten_17 (Flatten)	(None, 512)	0
batch_normalization_33 (Batch Normalization)	(None, 512)	2048
dense_60 (Dense)	(None, 128)	65664
dropout_15 (Dropout)	(None, 128)	0
batch_normalization_34 (Batch Normalization)	(None, 128)	512
dense_61 (Dense)	(None, 64)	8256
dropout_16 (Dropout)	(None, 64)	0
batch_normalization_35 (Batch Normalization)	(None, 64)	256
dense_62 (Dense)	(None, 2)	130
Total params: 14,791,554		
Trainable params: 14,790,146		
Non-trainable params: 1,408		

```
[218]: start=time.time()
history_vgg16_tuning = model.
    ↪ fit_generator(train_generator, epochs=10, validation_data=validation_generator)
runtime=time.time()-start+runtime
```

Epoch 1/10

66/66 [=====] - 99s 1s/step - loss: 0.1276 - accuracy: 0.9570 - val_loss: 0.1298 - val_accuracy: 0.9549

Epoch 2/10

66/66 [=====] - 73s 1s/step - loss: 0.1094 - accuracy: 0.9586 - val_loss: 0.1634 - val_accuracy: 0.9377

Epoch 3/10

66/66 [=====] - 74s 1s/step - loss: 0.0909 - accuracy: 0.9626 - val_loss: 0.0751 - val_accuracy: 0.9674

Epoch 4/10


```

66/66 [=====] - 74s 1s/step - loss: 0.0801 - accuracy:
0.9740 - val_loss: 0.0944 - val_accuracy: 0.9645
Epoch 5/10
66/66 [=====] - 74s 1s/step - loss: 0.0833 - accuracy:
0.9712 - val_loss: 0.1683 - val_accuracy: 0.9348
Epoch 6/10
66/66 [=====] - 74s 1s/step - loss: 0.0749 - accuracy:
0.9761 - val_loss: 0.1139 - val_accuracy: 0.9540
Epoch 7/10
66/66 [=====] - 74s 1s/step - loss: 0.0656 - accuracy:
0.9803 - val_loss: 0.0748 - val_accuracy: 0.9703
Epoch 8/10
66/66 [=====] - 74s 1s/step - loss: 0.0749 - accuracy:
0.9766 - val_loss: 0.0698 - val_accuracy: 0.9722
Epoch 9/10
66/66 [=====] - 74s 1s/step - loss: 0.0689 - accuracy:
0.9761 - val_loss: 0.0656 - val_accuracy: 0.9751
Epoch 10/10
66/66 [=====] - 75s 1s/step - loss: 0.0635 - accuracy:
0.9767 - val_loss: 0.0763 - val_accuracy: 0.9722

```

```

[219]: #saving the model and histories
#tf.keras.models.save_model(filepath=cwd, model=model)
model.save(cwd+'/assets/model_vgg16_tuning')
np.save(cwd+'/assets/history_vgg16_tuning.npy', history_vgg16_tuning.history)
#np.save(cwd+'/assets/history_xception_tuning.npy', history_xception_tuning.
→history)

```

```

INFO:tensorflow:Assets written to:
/home/jovyan/work/fc_project/assets/model_vgg16_tuning/assets

```

```

[220]: score = model.evaluate(test_generator, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

```

```

Test score: 0.17918924987316132
Test accuracy: 0.9294871687889099

```

```

[221]: results['model'].append('vgg16')
results['total_param'].append(14791146)
results['time'].append(runtime)
results['score'].append(score[1])
np.save(cwd+'/assets/results.npy', results)

```

```

[222]: STEP_SIZE_TEST=test_generator.n//test_generator.batch_size
test_generator.reset()
preds = model.predict(test_generator,
steps=STEP_SIZE_TEST,

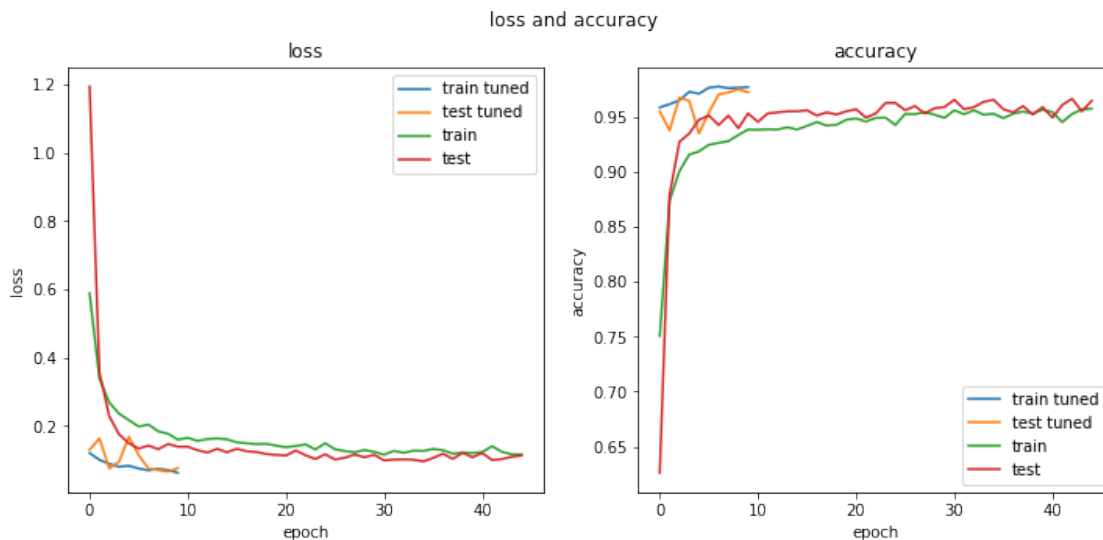
```

```
verbose=1)
```

624/624 [======] - 5s 8ms/step

```
[228]: fig, (ax1, ax2) = plt.subplots(1,2, figsize=(12,5))
ax1.plot(history_vgg16_tuning.history['loss'])
ax1.plot(history_vgg16_tuning.history['val_loss'])
ax1.plot(history_vgg16.history['loss'])
ax1.plot(history_vgg16.history['val_loss'])
ax1.set_title('loss')
ax1.set_ylabel('loss')
ax1.set_xlabel('epoch')
ax1.legend(['train tuned', 'test tuned', 'train', 'test'], loc='upper right')
ax2.plot(history_vgg16_tuning.history['accuracy'])
ax2.plot(history_vgg16_tuning.history['val_accuracy'])
ax2.plot(history_vgg16.history['accuracy'])
ax2.plot(history_vgg16.history['val_accuracy'])

ax2.set_title('accuracy')
ax2.set_ylabel('accuracy')
ax2.set_xlabel('epoch')
ax2.legend(['train tuned', 'test tuned', 'train', 'test'], loc='lower right')
fig.suptitle('loss and accuracy')
plt.show()
```



```
[238]: #scale the output of the model to [0,1]
#scaled_preds=(preds[:,1]-min(preds[:,1]))/(max(preds[:,1])-min(preds[:,1]))
pos = [i for i, j in zip(preds[:,1], test_generator.classes) if j == 1]
neg = [i for i, j in zip(preds[:,1], test_generator.classes) if j == 0]
```

```

with plt.xkcd():
    fig = plt.figure(figsize=(8, 4))

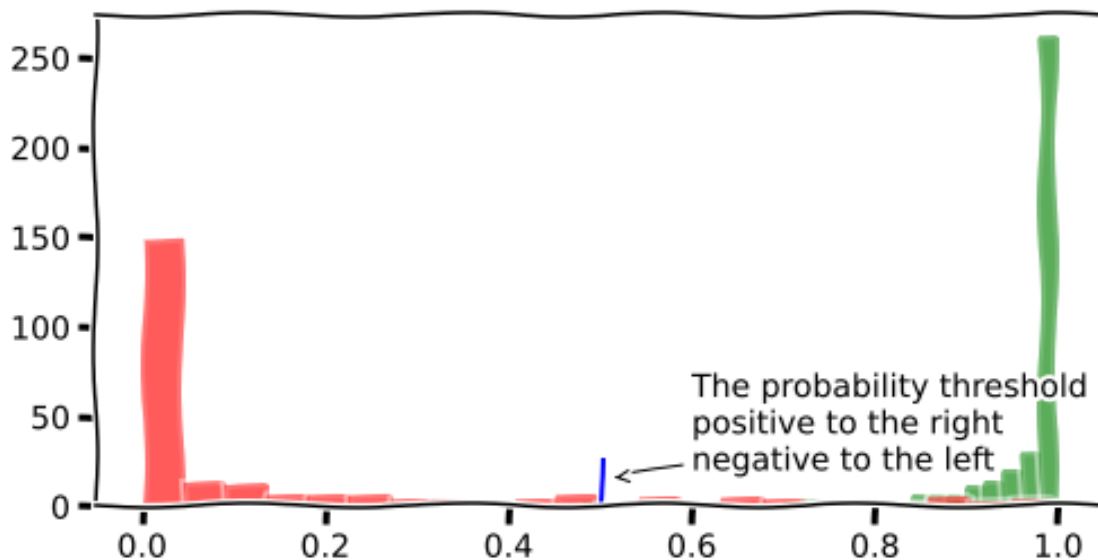
    sns.distplot(pos, hist = True, kde = False, color='g',
                 kde_kws = {'shade': True, 'linewidth': 3})

    sns.distplot(neg, hist = True, kde = False, color='r',
                 kde_kws = {'shade': True, 'linewidth': 3})

    plt.plot([0.5, 0.5], [0, 25], '-b')
    plt.annotate(
        'The probability threshold\npositive to the right\nnegative to the left',
        xy=(0.51, 15), arrowprops=dict(arrowstyle='->'), xytext=(0.6, 20))

plt.show()

```



```

[239]: confusion = confusion_matrix(test_generator.classes, np.where(preds[:,1]>0.
    ↪5,1,0), labels=[1, 0])
print(confusion)

```

```


[[373  17]
 [ 27 207]]

```

```

[240]: plot_confusion_matrix(cm=confusion, target_names = ['Space', 'Not Space'],
    ↪title = 'Confusion Matrix',normalize=False)

```




```
[241]: target_names = ['negative', 'positive']
print(classification_report(test_generator.classes, np.where(preds[:,1]>0.
↪5,1,0) , target_names=target_names))
```

	precision	recall	f1-score	support
negative	0.92	0.88	0.90	234
positive	0.93	0.96	0.94	390
accuracy			0.93	624
macro avg	0.93	0.92	0.92	624
weighted avg	0.93	0.93	0.93	624

```
[242]: fpr, tpr, _ = roc_curve(test_generator.classes, preds[:,1])
roc_auc = auc(fpr, tpr)
```

```
plt.figure()
lw = 2
plt.plot(fpr, tpr, color='darkorange',
lw=lw, label='ROC curve (area = %0.2f)' % roc_auc)
plt.plot([0, 1], [0, 1], color='navy', lw=lw, linestyle='--')
plt.xlim([-0.05, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic curve')
plt.legend(loc="lower right")
plt.show()
```



[]: